

Logistic Regression on Amazon Reviews Dataset

June 16, 2018

```
In [1]: import sqlite3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.plotly as py
import plotly.graph_objs as go
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.decomposition import TruncatedSVD
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/cross_validation.py:41: DeprecationWarning:
```

```
This module was deprecated in version 0.18 in favor of the model_selection module into which all
```

1 Loading and Sampling the dataset(100k data-points)

```
In [0]: final = pd.read_csv("final.csv")
```

```
final_data = final.sample(n = 100000)
```

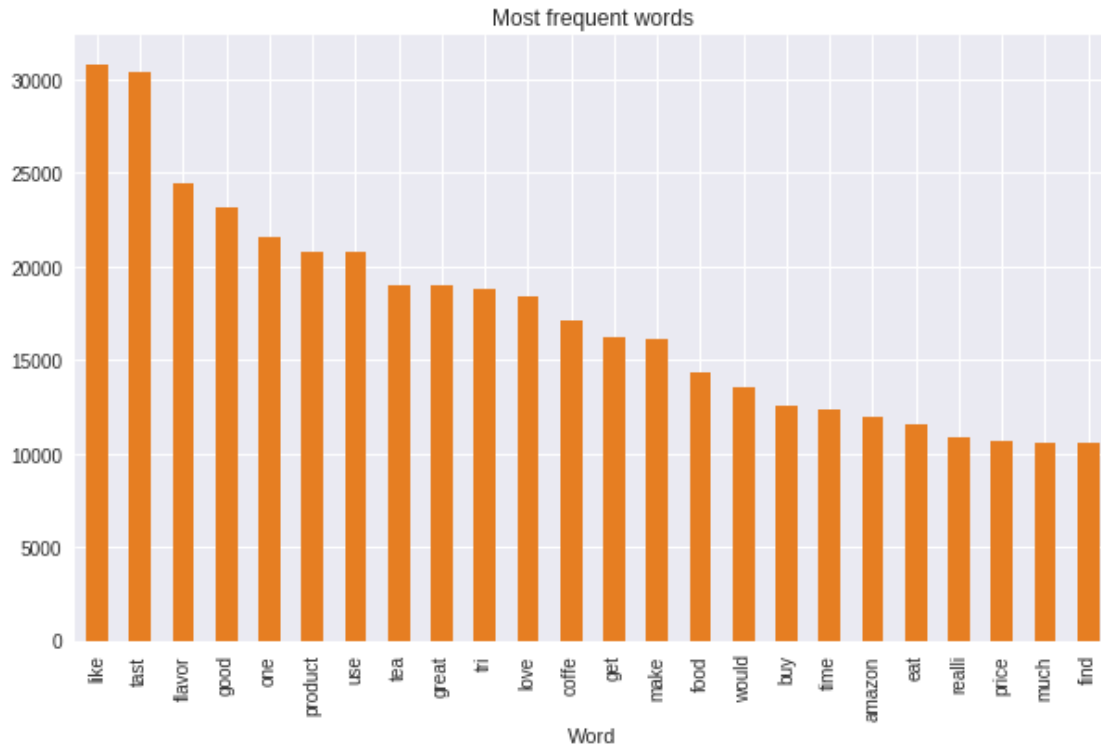


```
In [0]: print(word_plot)
```

	Word	Count
1	like	30800
2	tast	30378
3	flavor	24481
4	good	23182
5	one	21533
6	product	20781
7	use	20744
8	tea	19030
9	great	19021
10	tri	18764
11	love	18374
12	coffe	17075
13	get	16219
14	make	16114
15	food	14288
16	would	13576
17	buy	12549
18	time	12343
19	amazon	11952
20	eat	11543
21	realli	10904
22	price	10670
23	much	10596
24	find	10563

```
In [0]: word_plot.plot(kind='bar',x=word_plot['Word'],legend=False,title='Most frequent words')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fccc3566c50>
```



3.1 We can observe from the above plot that the most important words are:

- like
- taste
- flavour

4 Bag of words Vectorization

```
In [0]: count_vect = CountVectorizer() #in scikit-learn
X_train1 = count_vect.fit_transform(X_train)
X_test1 = count_vect.transform(X_test)
```

```
In [0]: #Standardization
from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train1 = sc.fit_transform(X_train1)
X_test1 = sc.transform(X_test1)
```

/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:475: DataConversionWarning:

Data with input dtype int64 was converted to float64 by StandardScaler.

5 GridSearch Cross-Validation

```
In [0]: my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train1)]

        param_grid = {'C': [0.001, 0.01, 0.1, 1, 10]}

        grid = GridSearchCV(LogisticRegression(), param_grid, cv=my_cv)
        grid.fit(X_train1, y_train)

Out[0]: GridSearchCV(cv=[(array([ 0,    1, ..., 6368, 6369]), array([ 6370, 6371, ..., 1273
        error_score='raise',
        estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
        intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
        verbose=0, warm_start=False),
        fit_params=None, iid=True, n_jobs=1,
        param_grid={'C': [0.001, 0.01, 0.1, 1, 10]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
        scoring=None, verbose=0)

In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
        print("Best parameters: ", grid.best_params_)

Best cross-validation score: 0.899
Best parameters: {'C': 0.001}
```

6 RandomizedSearch Cross-Validation

```
In [0]: from sklearn.model_selection import RandomizedSearchCV
        from scipy import stats

        my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train1)]

        param_grid = {"C": stats.uniform(0.001, 10)}

        grid = RandomizedSearchCV(LogisticRegression(), param_grid, cv=my_cv, n_iter = 10)
        grid.fit(X_train1, y_train)

Out[0]: RandomizedSearchCV(cv=[(array([ 0,    1, ..., 6368, 6369]), array([ 6370, 6371, ...
        error_score='raise',
        estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
        intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
        verbose=0, warm_start=False),
        fit_params=None, iid=True, n_iter=10, n_jobs=1,
```

```

param_distributions={'C': <scipy.stats._distn_infrastructure.rv_frozen object>
pre_dispatch='2*n_jobs', random_state=None, refit=True,
return_train_score='warn', scoring=None, verbose=0)

```

```

In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
        print("Best parameters: ", grid.best_params_)

```

Best cross-validation score: 0.869

Best parameters: {'C': 0.6203520096925298}

7 Applying Logistic Regression

```

In [0]: logreg = LogisticRegression(C = 0.001).fit(X_train1, y_train)

```

```

print("Training set score: {:.3f}".format(logreg.score(X_train1, y_train)))
print("Test set score: {:.3f}".format(logreg.score(X_test1, y_test)))

```

Training set score: 0.970

Test set score: 0.901

8 Generating a Confusion matrix

```

In [0]: y_pred = logreg.predict(X_test1)

```

```

log_cfm = confusion_matrix(y_test, y_pred)

```

```

In [0]: import itertools

```

```

plt.imshow(log_cfm, interpolation='nearest')

```

```


for i, j in itertools.product(range(log_cfm.shape[0]), range(log_cfm.shape[1])):
    plt.text(j, i, log_cfm[i, j],
             horizontalalignment="center",
             color="red")

```

```

plt.ylabel('True label (Recall)')
plt.xlabel('Predicted label (Precision)')
plt.title('Logistic Reg | Confusion Matrix')
plt.colorbar();

```



9 Comparing 'L1' and 'L2' regularization techniques

```
In [0]: for i, C in enumerate((0.001, 0.01, 0.1, 1)):
        # turn down tolerance for short training time
        clf_l1_LR = LogisticRegression(C=C, penalty='l1', tol=0.01)
        clf_l2_LR = LogisticRegression(C=C, penalty='l2', tol=0.01)
        clf_l1_LR.fit(X_train1, y_train)
        clf_l2_LR.fit(X_train1, y_train)

        coef_l1_LR = clf_l1_LR.coef_.ravel()
        coef_l2_LR = clf_l2_LR.coef_.ravel()

        # coef_l1_LR contains zeros due to the
        # L1 sparsity inducing norm

        sparsity_l1_LR = np.mean(coef_l1_LR == 0) * 100
        sparsity_l2_LR = np.mean(coef_l2_LR == 0) * 100
```

```

print("C=%.3f" % C)
print("Sparsity with L1 penalty: %.3f%" % sparsity_l1_LR)
print("score with L1 penalty: %.4f" % clf_l1_LR.score(X_train1, y_train))
print("Sparsity with L2 penalty: %.3f%" % sparsity_l2_LR)
print("score with L2 penalty: %.4f" % clf_l2_LR.score(X_train1, y_train))

```

```

C=0.001
Sparsity with L1 penalty: 99.759%
score with L1 penalty: 0.8743
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.9695
C=0.010
Sparsity with L1 penalty: 88.691%
score with L1 penalty: 0.9404
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.9805
C=0.100
Sparsity with L1 penalty: 65.301%
score with L1 penalty: 0.9767
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.9856
C=1.000
Sparsity with L1 penalty: 52.631%
score with L1 penalty: 0.9848
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.9872

```

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	99.7 %	0.87
L2-regularizer	0.001	0 %	0.96
---	---	---	---
L1-regularizer	0.01	88.6%	0.94
L2-regularizer	0.01	0 %	0.98
---	---	---	---
L1-regularizer	0.1	65.3%	0.97
L2-regularizer	0.1	0 %	0.98
---	---	---	---
L1-regularizer	1.0	52.6%	0.98
L2-regularizer	1.0	0 %	0.98

10 TF-idf Vectorization

```

In [0]: #TF-IDF
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))

```



```
X_train2 = tf_idf_vect.fit_transform(X_train)
X_test2 = tf_idf_vect.transform(X_test)
```

```
In [0]: #Standardization
from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train2 = sc.fit_transform(X_train2)
X_test2 = sc.transform(X_test2)
```

11 Applying GridSearch CV

```
In [0]: my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train2)]

param_grid = {'C': [0.001, 0.01, 0.1, 1, 10]}

grid = GridSearchCV(LogisticRegression(), param_grid, cv=my_cv)
grid.fit(X_train2, y_train)
```

```
Out[0]: GridSearchCV(cv=[(array([ 0, 1, ..., 6368, 6369]), array([ 6370, 6371, ..., 1273
error_score='raise',
estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False),
fit_params=None, iid=True, n_jobs=1,
param_grid={'C': [0.001, 0.01, 0.1, 1, 10]}],
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=0)
```

```
In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
print("Best parameters: ", grid.best_params_)
```

```
Best cross-validation score: 0.857
Best parameters: {'C': 10}
```

12 Applying RandomizedSearch CV

```
In [0]: from sklearn.model_selection import RandomizedSearchCV
from scipy import stats
```

```
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train2)]

param_grid = {"C": stats.uniform(0.001, 10)}
```

```
grid = RandomizedSearchCV(LogisticRegression(), param_grid, cv=my_cv, n_iter = 10)
grid.fit(X_train2, y_train)
```

```
Out[0]: RandomizedSearchCV(cv=[(array([ 0, 1, ..., 6368, 6369]), array([ 6370, 6371, ...
      error_score='raise',
      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
      intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
      penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
      verbose=0, warm_start=False),
      fit_params=None, iid=True, n_iter=1, n_jobs=1,
      param_distributions={'C': <scipy.stats._distn_infrastructure.rv_frozen object>},
      pre_dispatch='2*n_jobs', random_state=None, refit=True,
      return_train_score='warn', scoring=None, verbose=0)
```

```
In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
      print("Best parameters: ", grid.best_params_)
```

```
Best cross-validation score: 0.857
Best parameters: {'C': 1.9773164474979121}
```

```
In [0]: logreg = LogisticRegression(C = 10.0).fit(X_train2, y_train)
```

```
      print("Training set score: {:.3f}".format(logreg.score(X_train2, y_train)))
      print("Test set score: {:.3f}".format(logreg.score(X_test2, y_test)))
```

```
Training set score: 1.000
Test set score: 0.864
```

13 Generating Confusion matrix

```
In [0]: y_pred = logreg.predict(X_test2)
```

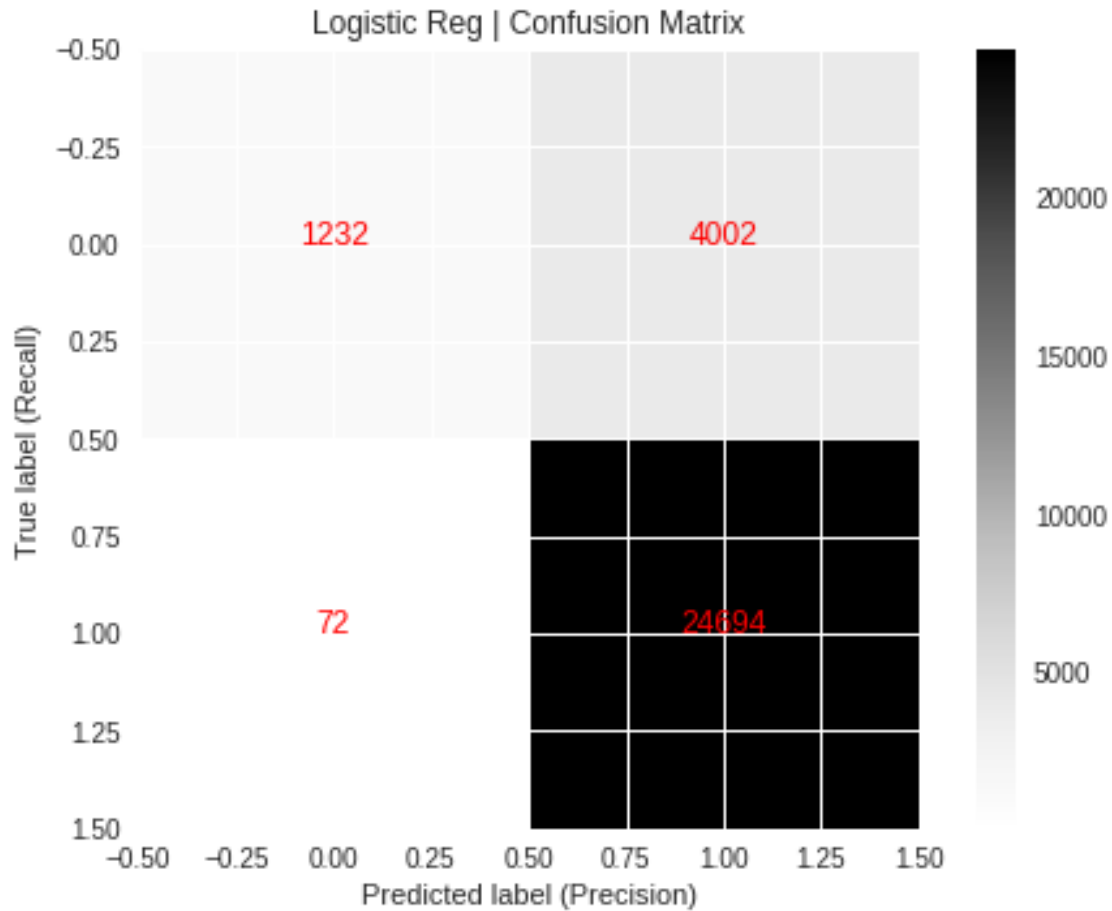
```
      log_cfm = confusion_matrix(y_test, y_pred)
```

```
In [0]: import itertools
```

```
      plt.imshow(log_cfm, interpolation='nearest')
```

```
      for i, j in itertools.product(range(log_cfm.shape[0]), range(log_cfm.shape[1])):
          plt.text(j, i, log_cfm[i, j],
                  horizontalalignment="center",
                  color="red")
```

```
      plt.ylabel('True label (Recall)')
      plt.xlabel('Predicted label (Precision)')
      plt.title('Logistic Reg | Confusion Matrix')
      plt.colorbar();
```



```
In [0]: for i, C in enumerate((0.001, 0.01, 0.1, 1)):
        # turn down tolerance for short training time
        clf_l1_LR = LogisticRegression(C=C, penalty='l1', tol=0.01)
        clf_l2_LR = LogisticRegression(C=C, penalty='l2', tol=0.01)
        clf_l1_LR.fit(X_train2, y_train)
        clf_l2_LR.fit(X_train2, y_train)

        coef_l1_LR = clf_l1_LR.coef_.ravel()
        coef_l2_LR = clf_l2_LR.coef_.ravel()

        # coef_l1_LR contains zeros due to the
        # L1 sparsity inducing norm

        sparsity_l1_LR = np.mean(coef_l1_LR == 0) * 100
        sparsity_l2_LR = np.mean(coef_l2_LR == 0) * 100

        print("C=%.3f" % C)
        print("Sparsity with L1 penalty: %.3f%%" % sparsity_l1_LR)
```

```

print("score with L1 penalty: %.4f" % clf_l1_LR.score(X_train2, y_train))
print("Sparsity with L2 penalty: %.3f%" % sparsity_l2_LR)
print("score with L2 penalty: %.4f" % clf_l2_LR.score(X_train2, y_train))

```

```

C=0.001
Sparsity with L1 penalty: 99.992%
score with L1 penalty: 0.8777
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 1.0000
C=0.010
Sparsity with L1 penalty: 97.911%
score with L1 penalty: 0.9996
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 1.0000
C=0.100
Sparsity with L1 penalty: 94.378%
score with L1 penalty: 1.0000
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 1.0000
C=1.000
Sparsity with L1 penalty: 85.909%
score with L1 penalty: 1.0000
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 1.0000

```

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	99.9 %	0.87
L2-regularizer	0.001	0 %	1.0
---	---	---	---
L1-regularizer	0.01	97.9%	0.99
L2-regularizer	0.01	0 %	1.0
---	---	---	---
L1-regularizer	0.1	94.3%	1.0
L2-regularizer	0.1	0 %	1.0
---	---	---	---
L1-regularizer	1.0	85.9%	1.0
L2-regularizer	1.0	0 %	1.0

14 Word2Vec Vectorization

```

In [5]: import nltk
        nltk.download('stopwords')
        import re
        # Tutorial about Python regular expressions: https://pymotw.com/2/re/
        import string

```

```

from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
def cleanpunc(sentence): #function to clean the word of any punctuation or special cha
    cleaned = re.sub(r'[?|!|\'|\"|#]',r'',sentence)
    cleaned = re.sub(r'[.,|)|(|\\|/]',r' ',cleaned)
    return cleaned
print(stop)
print('*****')
print(sno.stem('tasty'))

```

[nltk_data] Downloading package stopwords to /content/nltk_data...

[nltk_data] Package stopwords is already up-to-date!

{"that'll", 'his', 've', 'we', 'how', 'about', 'some', 'didn', 'an', 'for', 'y', 'down', "could

tasti

14.1 Training Word2Vec model using own text corpus

```

In [0]: import gensim
        i=0
        train_sent=[]
        for sent in X_train:
            filtered_sentence=[]
            sent=cleanhtml(sent)
            for w in sent.split():
                for cleaned_words in cleanpunc(w).split():
                    if(cleaned_words.isalpha()):
                        filtered_sentence.append(cleaned_words.lower())
                    else:
                        continue
            train_sent.append(filtered_sentence)

```

```

In [0]: test_sent=[]
        for sent in X_test:
            filtered_sentence=[]
            sent=cleanhtml(sent)
            for w in sent.split():

```

```

        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    test_sent.append(filtered_sentence)

```

```

In [0]: from gensim.models import Word2Vec
        from gensim.models import KeyedVectors

```

```

w2v_model=gensim.models.Word2Vec(train_sent,min_count=5,size=50, workers=4)

```

15 Applying Average Word2vec

```

In [0]: #AVG-W2V
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in train_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

70000

50

```

In [0]: #AVG-W2V
sent_vectors2 = []; # the avg-w2v for each sentence/review is stored in this list
for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words

```

```

        sent_vectors2.append(sent_vec)
print(len(sent_vectors2))
print(len(sent_vectors2[0]))

```

30000

50

```

In [0]: #Standardization
from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train3 = sc.fit_transform(sent_vectors)
X_test3 = sc.transform(sent_vectors2)

```

16 Applying GridSearch CV

```

In [0]: my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train3)]

param_grid = {'C': [0.001, 0.01, 0.1, 1, 10]}

grid = GridSearchCV(LogisticRegression(), param_grid, cv=my_cv)
grid.fit(X_train3, y_train)

```

```

Out[0]: GridSearchCV(cv=[(array([ 0, 1, ..., 6368, 6369]), array([ 6370, 6371, ..., 1273
error_score='raise',
estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False),
fit_params=None, iid=True, n_jobs=1,
param_grid={'C': [0.001, 0.01, 0.1, 1, 10]}},
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=0)

```

```

In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
print("Best parameters: ", grid.best_params_)

```

Best cross-validation score: 0.887

Best parameters: {'C': 0.1}

17 Applying Randomized CV

```

In [0]: from sklearn.model_selection import RandomizedSearchCV
from scipy import stats

```

```

my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train3)]

```

```
param_grid = {"C": stats.uniform(0.001, 10)}
```

```
grid = RandomizedSearchCV(LogisticRegression(), param_grid, cv=my_cv, n_iter = 10)
grid.fit(X_train3, y_train)
```

```
Out[0]: RandomizedSearchCV(cv=[(array([ 0, 1, ..., 6368, 6369]), array([ 6370, 6371, ...
      error_score='raise',
      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
      intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
      penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
      verbose=0, warm_start=False),
      fit_params=None, iid=True, n_iter=1, n_jobs=1,
      param_distributions={'C': <scipy.stats._distn_infrastructure.rv_frozen object>},
      pre_dispatch='2*n_jobs', random_state=None, refit=True,
      return_train_score='warn', scoring=None, verbose=0)
```

```
In [0]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
        print("Best parameters: ", grid.best_params_)
```

Best cross-validation score: 0.887

Best parameters: {'C': 5.16712437106295}

```
In [0]: logreg = LogisticRegression(C = 0.1).fit(X_train3, y_train)
```

```
print("Training set score: {:.3f}".format(logreg.score(X_train3, y_train)))
print("Test set score: {:.3f}".format(logreg.score(X_test3, y_test)))
```

Training set score: 0.891

Test set score: 0.877

18 Generating Confusion matrix

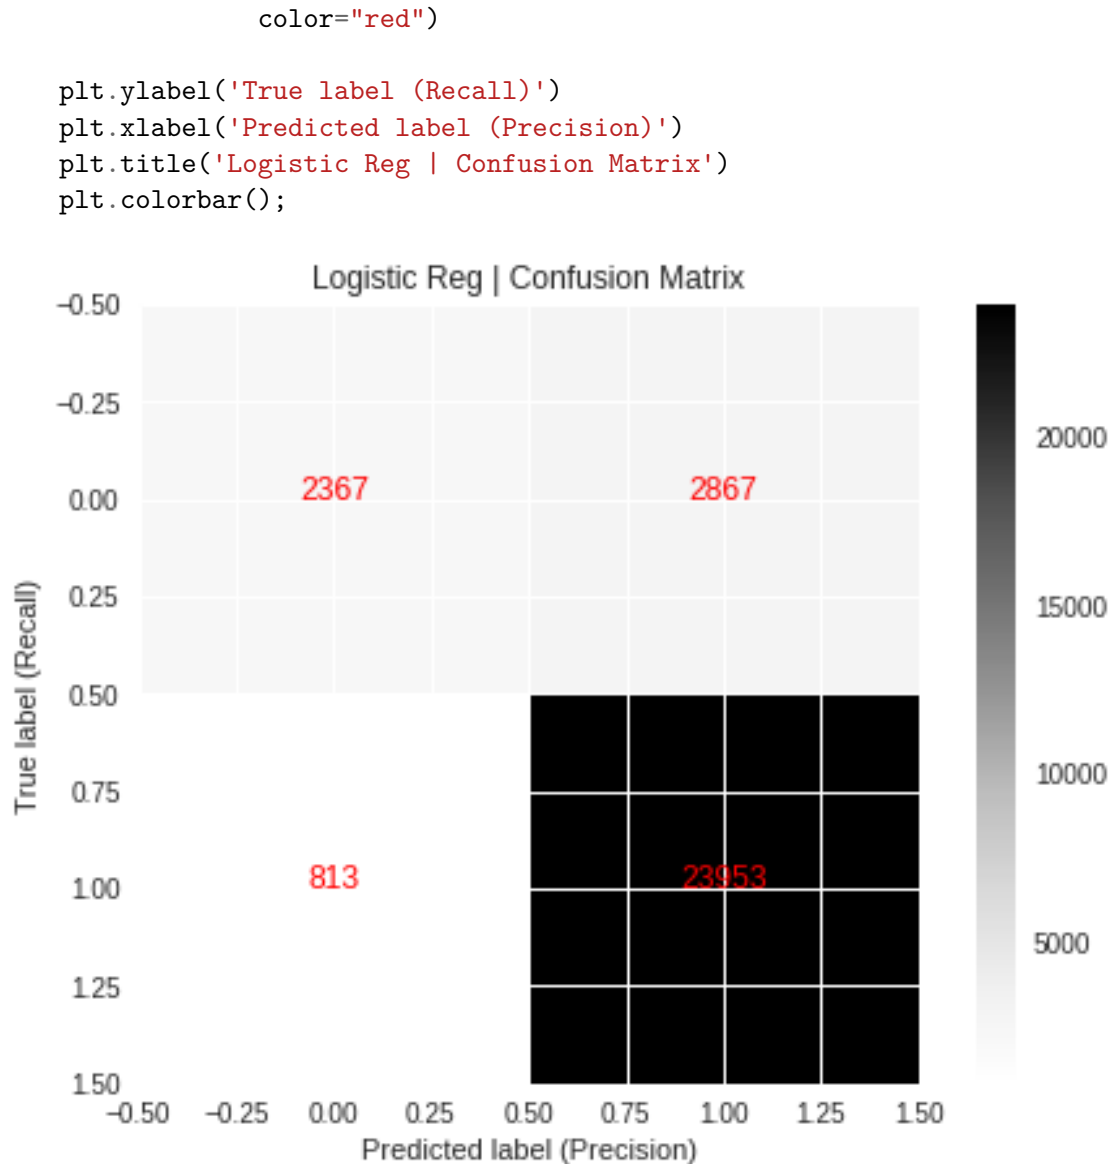
```
In [0]: y_pred = logreg.predict(X_test3)
```

```
log_cfm = confusion_matrix(y_test, y_pred)
```

```
In [0]: import itertools
```

```
plt.imshow(log_cfm, interpolation='nearest')
```

```
for i, j in itertools.product(range(log_cfm.shape[0]), range(log_cfm.shape[1])):
    plt.text(j, i, log_cfm[i, j],
             horizontalalignment="center",
```

```
In [0]: for i, C in enumerate((0.001, 0.01, 0.1, 1)):
        # turn down tolerance for short training time
        clf_l1_LR = LogisticRegression(C=C, penalty='l1', tol=0.01)
        clf_l2_LR = LogisticRegression(C=C, penalty='l2', tol=0.01)
        clf_l1_LR.fit(X_train3, y_train)
        clf_l2_LR.fit(X_train3, y_train)

        coef_l1_LR = clf_l1_LR.coef_.ravel()
        coef_l2_LR = clf_l2_LR.coef_.ravel()

        # coef_l1_LR contains zeros due to the
```

```

# L1 sparsity inducing norm

sparsity_l1_LR = np.mean(coef_l1_LR == 0) * 100
sparsity_l2_LR = np.mean(coef_l2_LR == 0) * 100

print("C=%.3f" % C)
print("Sparsity with L1 penalty: %.3f%%" % sparsity_l1_LR)
print("score with L1 penalty: %.4f" % clf_l1_LR.score(X_train3, y_train))
print("Sparsity with L2 penalty: %.3f%%" % sparsity_l2_LR)
print("score with L2 penalty: %.4f" % clf_l2_LR.score(X_train3, y_train))

C=0.001
Sparsity with L1 penalty: 60.000%
score with L1 penalty: 0.8749
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8885
C=0.010
Sparsity with L1 penalty: 14.000%
score with L1 penalty: 0.8900
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8911
C=0.100
Sparsity with L1 penalty: 0.000%
score with L1 penalty: 0.8910
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8913
C=1.000
Sparsity with L1 penalty: 0.000%
score with L1 penalty: 0.8909
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8913

```

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	60 %	0.87
L2-regularizer	0.001	0 %	0.88
---	---	---	---
L1-regularizer	0.01	14 %	0.89
L2-regularizer	0.01	0 %	0.89
---	---	---	---
L1-regularizer	0.1	0 %	0.89
L2-regularizer	0.1	0 %	0.89
---	---	---	---
L1-regularizer	1.0	0%	0.89
L2-regularizer	1.0	0 %	0.89

19 Word2Vec Tf-idf Vectorization

In [0]: *#TF-IDF*

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf = tf_idf_vect.fit_transform(X_train)
tfidf_feat = tf_idf_vect.get_feature_names()
# tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in train_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf

        except:
            pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    train_vectors.append(sent_vec)
    row += 1
```

In [0]: *#TF-IDF*

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf = tf_idf_vect.fit_transform(X_test)
tfidf_feat = tf_idf_vect.get_feature_names()
# tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
```

```

        vec = w2v_model2.wv[word]
        # obtain the tf_idf of a word in a sentence/review
        tfidf = final_tf_idf[row, tfidf_feat.index(word)]
        sent_vec += (vec * tfidf)
        weight_sum += tfidf

    except:
        pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    test_vectors.append(sent_vec)
    row += 1

```

```

In [0]: #Standardization
        from sklearn.preprocessing import StandardScaler
        sc= StandardScaler(with_mean=False)
        X_train4 = sc.fit_transform(train_vectors)
        X_test4 = sc.transform(test_vectors)

```

20 Applying GridSearch CV

```

In [16]: my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train4)]

```

```

param_grid = {'C': [0.001, 0.01, 0.1, 1, 10]}

grid = GridSearchCV(LogisticRegression(), param_grid, cv=my_cv)
grid.fit(X_train4, y_train)

```

```

Out[16]: GridSearchCV(cv=[(array([ 0,  1, ..., 638, 639]), array([ 640,  641, ..., 1274, 1275]),
      error_score='raise',
      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
      intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
      penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
      verbose=0, warm_start=False),
      fit_params=None, iid=True, n_jobs=1,
      param_grid={'C': [0.001, 0.01, 0.1, 1, 10]},
      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
      scoring=None, verbose=0)

```

```

In [17]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
        print("Best parameters: ", grid.best_params_)

```

```

Best cross-validation score: 0.850
Best parameters: {'C': 10}

```

21 Applying RandomizedSearch CV

```
In [18]: from sklearn.model_selection import RandomizedSearchCV
         from scipy import stats
```

```
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train4)]

param_grid = {"C": stats.uniform(0.001, 10)}
```

```
grid = RandomizedSearchCV(LogisticRegression(), param_grid, cv=my_cv, n_iter = 10)
grid.fit(X_train4, y_train)
```

```
Out[18]: RandomizedSearchCV(cv=[(array([ 0,  1, ..., 638, 639]), array([ 640,  641, ..., 127
      error_score='raise',
      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
      intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
      penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
      verbose=0, warm_start=False),
      fit_params=None, iid=True, n_iter=1, n_jobs=1,
      param_distributions={'C': <scipy.stats._distn_infrastructure.rv_frozen object>},
      pre_dispatch='2*n_jobs', random_state=None, refit=True,
      return_train_score='warn', scoring=None, verbose=0)
```

```
In [19]: print("Best cross-validation score: {:.3f}".format(grid.best_score_))
         print("Best parameters: ", grid.best_params_)
```

```
Best cross-validation score: 0.850
Best parameters:  {'C': 9.080485889967418}
```

```
In [20]: logreg = LogisticRegression(C = 10).fit(X_train4, y_train)
```

```
print("Training set score: {:.3f}".format(logreg.score(X_train4, y_train)))
print("Test set score: {:.3f}".format(logreg.score(X_test4, y_test)))
```

```
Training set score: 0.858
Test set score: 0.825
```

22 Generating Confusion matrix

```
In [0]: y_pred = logreg.predict(X_test4)
```

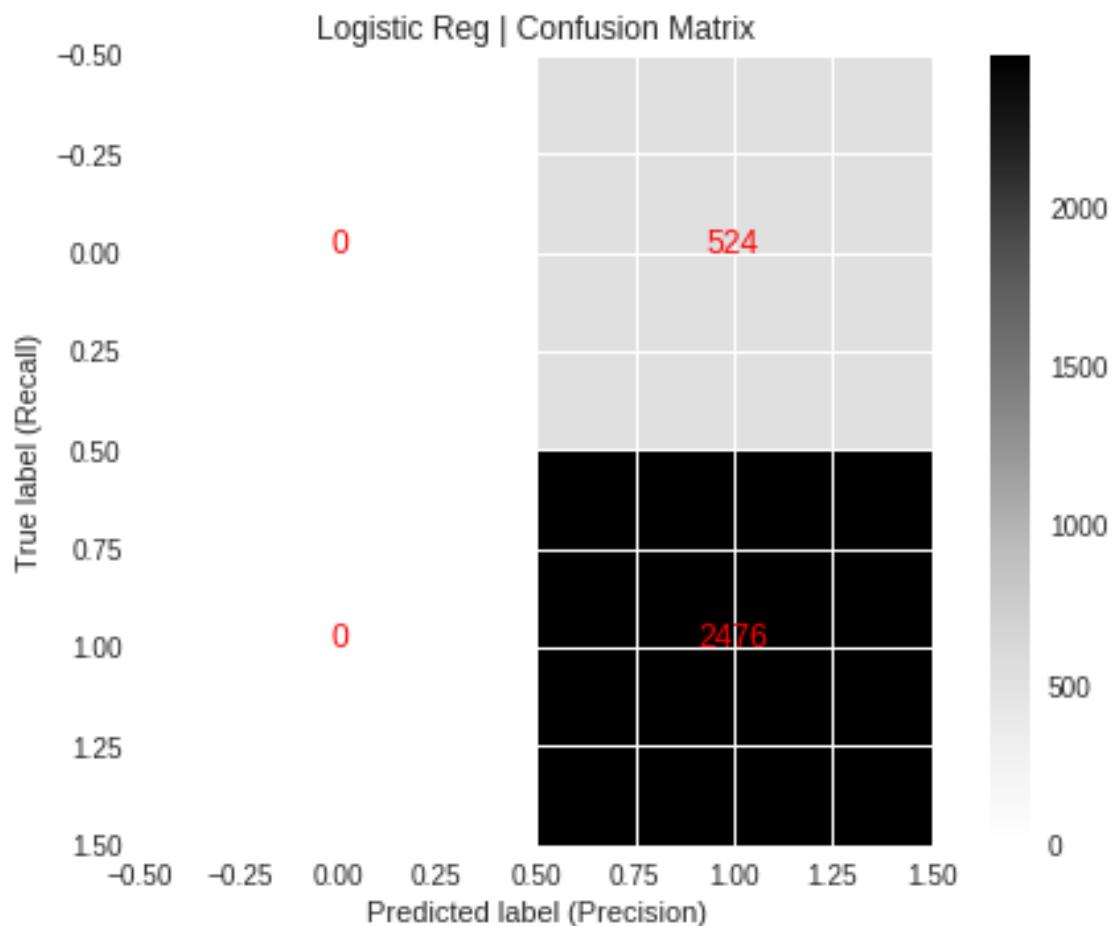
```
log_cfm = confusion_matrix(y_test, y_pred)
```

```
In [22]: import itertools
```

```
plt.imshow(log_cfm, interpolation='nearest')

for i, j in itertools.product(range(log_cfm.shape[0]), range(log_cfm.shape[1])):
    plt.text(j, i, log_cfm[i, j],
             horizontalalignment="center",
             color="red")

plt.ylabel('True label (Recall)')
plt.xlabel('Predicted label (Precision)')
plt.title('Logistic Reg | Confusion Matrix')
plt.colorbar();
```



```
In [23]: for i, C in enumerate((0.001, 0.01, 0.1, 1, 10.0)):
    # turn down tolerance for short training time
    clf_l1_LR = LogisticRegression(C=C, penalty='l1', tol=0.01)
    clf_l2_LR = LogisticRegression(C=C, penalty='l2', tol=0.01)
```

```

clf_l1_LR.fit(X_train4, y_train)
clf_l2_LR.fit(X_train4, y_train)

coef_l1_LR = clf_l1_LR.coef_.ravel()
coef_l2_LR = clf_l2_LR.coef_.ravel()

# coef_l1_LR contains zeros due to the
# L1 sparsity inducing norm

sparsity_l1_LR = np.mean(coef_l1_LR == 0) * 100
sparsity_l2_LR = np.mean(coef_l2_LR == 0) * 100

print("C=%.3f" % C)
print("Sparsity with L1 penalty: %.3f%%" % sparsity_l1_LR)
print("score with L1 penalty: %.4f" % clf_l1_LR.score(X_train4, y_train))
print("Sparsity with L2 penalty: %.3f%%" % sparsity_l2_LR)
print("score with L2 penalty: %.4f" % clf_l2_LR.score(X_train4, y_train))

```

```

C=0.001
Sparsity with L1 penalty: 90.000%
score with L1 penalty: 0.8514
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8516
C=0.010
Sparsity with L1 penalty: 80.000%
score with L1 penalty: 0.8514
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8514
C=0.100
Sparsity with L1 penalty: 28.000%
score with L1 penalty: 0.8519
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8520
C=1.000
Sparsity with L1 penalty: 0.000%
score with L1 penalty: 0.8516
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8531
C=10.000
Sparsity with L1 penalty: 0.000%
score with L1 penalty: 0.8521
Sparsity with L2 penalty: 0.000%
score with L2 penalty: 0.8577

```

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	90%	0.85

	Hyper-parameter	Sparsity	Score
L2-regularizer	0.001	0 %	0.85
---	---	---	---
L1-regularizer	0.01	80%	0.85
L2-regularizer	0.01	0 %	0.85
---	---	---	---
L1-regularizer	0.1	28 %	0.85
L2-regularizer	0.1	0 %	0.85
---	---	---	---
L1-regularizer	1.0	0%	0.85
L2-regularizer	1.0	0 %	0.85
---	---	---	---
L1-regularizer	10.0	0%	0.85
L2-regularizer	10.0	0%	0.85

23 Conclusion

23.1 The most important words are:

-

23.2 tast

-

23.3 like

-

23.4 product

23.5 Bag of Words:

-

23.5.1 GridSearch CV

- Best Parameter C = 0.001
- CV Score = 0.854
-

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	99.7 %	0.87
L2-regularizer	0.001	0 %	0.96
---	---	---	---
L1-regularizer	0.01	88.6%	0.94
L2-regularizer	0.01	0 %	0.98
---	---	---	---
L1-regularizer	0.1	65.3%	0.97
L2-regularizer	0.1	0 %	0.98
---	---	---	---
L1-regularizer	1.0	52.6%	0.98
L2-regularizer	1.0	0 %	0.98

23.5.2 RandomizedSearch CV

- Best Parameter C = 0.54
- CV Score = 0.796
-

23.6 Test set score = 0.90

-

23.7 Train set score = 0.97

23.8 Tf-IDF:

-

23.8.1 GridSearch CV

- Best Parameter C = 0.001
- CV Score = 0.824
-

23.8.2 RandomizedSearch CV

- Best Parameter C = 1.79
- CV Score = 0.823
-

23.9 Test set score = 0.86

-

23.10 Train set score = 1.0

23.11 Average Word2vec:

-

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	99.9 %	0.87
L2-regularizer	0.001	0 %	1.0
---	---	---	---
L1-regularizer	0.01	97.9 %	0.99
L2-regularizer	0.01	0 %	1.0
---	---	---	---
L1-regularizer	0.1	94.3 %	1.0
L2-regularizer	0.1	0 %	1.0
---	---	---	---
L1-regularizer	1.0	85.9 %	1.0
L2-regularizer	1.0	0 %	1.0

23.11.1 GridSearch CV

- Best Parameter C = 0.001
- CV Score = 0.832
-

23.11.2 RandomizedSearch CV

- Best Parameter C = 4.27
- CV Score = 0.831
-

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	60 %	0.87
L2-regularizer	0.001	0 %	0.88
---	---	---	---
L1-regularizer	0.01	14 %	0.89
L2-regularizer	0.01	0 %	0.89
---	---	---	---
L1-regularizer	0.1	0 %	0.89
L2-regularizer	0.1	0 %	0.89
---	---	---	---
L1-regularizer	1.0	0 %	0.89
L2-regularizer	1.0	0 %	0.89

23.12 Test set score = 0.87

-

23.13 Train set score = 0.89

23.14 Word2vec-tfidf:

-

23.14.1 GridSearch CV

- Best Parameter C = 10.0
- CV Score = 0.738
-

23.14.2 RandomizedSearch CV

- Best Parameter C = 8.42
- CV Score = 0.737
-

	Hyper-parameter	Sparsity	Score
L1-regularizer	0.001	90 %	0.85
L2-regularizer	0.001	0 %	0.85
---	---	---	---
L1-regularizer	0.01	80 %	0.85
L2-regularizer	0.01	0 %	0.85
---	---	---	---
L1-regularizer	0.1	28 %	0.85
L2-regularizer	0.1	0 %	0.85
---	---	---	---
L1-regularizer	1.0	0 %	0.85
L2-regularizer	1.0	0 %	0.85
---	---	---	---
L1-regularizer	10.0	0 %	0.85
L2-regularizer	10.0	0 %	0.85

23.15 Test set score = 0.82

-

23.16 Train set score = 0.85